The background of the slide is a photograph of a wind farm. Several wind turbines are visible, receding into the distance. The sky is a clear, deep blue. In the foreground, there are some dry, yellowish-brown plants. Overlaid on this image is a white geometric pattern consisting of several interconnected lines forming a series of triangles. Four white dots are placed at the vertices of these triangles: one at the top right, one at the bottom left, one in the middle left, and one at the bottom right. A large, semi-transparent green rectangle is positioned on the left side of the slide, containing the title text.

# Bayesian Machine Learning on determining the the power factor

alliander

# What is there to be learned?

- Classically Energy flow: Producer → Consumer
  - Easy to maintain balance between demand and supply
  - General: Single direction of power flow
- With decentralised power generation
  - Power flow in Middelspannings (MS) net becomes muddled
  - Hard to keep up with supply, demand, and capacity



→ Big bottleneck in the energy transition
- Current measurements insufficient to understand
  - Most MS substations only record absolute current (I) and voltages (U)
  - Look for a machine learning solution to learn the active power (P)

# What is there to be learned?

## Quick recap

### Alternating current grid:

- $S = U \times I \times \sqrt{3}$ 
  - Apparent power (S)
  - Theoretical maximum
- $S = \sqrt{P^2 + Q^2}$ 
  - Apparent power (S)
  - Hypotenuse of the active power (P) and the reactive power (Q).
  - P is the actually used/consumed power
  - Q is the non-used power → wasted
- $P/S = \cos(\phi)$ 
  - P/S is the power factor [-1,1]
  - P and Q are determined by the phase angle ( $\phi$ )
- If we can learn P from available data we are able to calculate the power factor at MS substations



# Quick overview of bayesian machine learning

Why Bayesian machine learning?

alltander

- It allows for statistical models that can be updated with prior knowledge expressed as distributions
  - Observational data is combined with the prior and normalised to give a posterior prediction
- Great! Perhaps we can encode common constraints of the grid in the prior to help in predicting power flow:
  - E.g. Kirchhoff's current law (  $\sum(P) = 0$  )
  - $P \leq S$  (realistically always  $P < S$ )

$$\underbrace{P(A|B)}_{\text{posterior}} = \underbrace{P(A)}_{\text{prior}} \times \frac{\underbrace{P(B|A)}_{\text{likelihood}}}{\underbrace{P(B)}_{\text{marginal}}}$$

# Quick overview of bayesian machine learning

## Gaussian Processes

alltander

- I decided to try apply Gaussian Processes (GP) as the Bayesian machine learning method
  - Specifically GP regression
- GP Regression (GPR)
  - Nonparametric: Not limited by a function form
  - GPR → distribution over all admissible functions fitting the data
  - Can be described as an infinite-dimensional multivariate gaussian distribution

$$f(x) \sim GP(m(x), k(x, x'))$$

- Gaussian prior
  - Mean function:  $m(x)$
  - Covariance function:  $k(x, x')$

$$\begin{bmatrix} y \\ f_* \end{bmatrix} \sim \mathcal{N} \left( \begin{bmatrix} \mu \\ \mu_* \end{bmatrix}, \begin{bmatrix} K(X, X) + \sigma_n^2 I & K(X, X_*) \\ K(X_*, X) & K(X_*, X_*) \end{bmatrix} \right)$$

- Training data
  - Gets applied to the covariance matrix to create the kernel matrix
  - Does not require a lot of data to train
- GPs allow for:
  - Interpolation of observations (useful for holes in the data)
  - Probabilistic predictions (confidence intervals)
- GP weaknesses
  - Computationally costly for high dimension feature spaces

# What I ended up doing

Three uneven topics

alllander

- Main focus on Bayesian machine learning / Gaussian processes
- However two (sort of) separate ideas

## 1 Look at the possible use of flow networks for energy grids

- One master thesis seemed promising
  - “*Network flow models for power grids*”, Francesca Wegner, 2014
- Required knowledge of wider energy flow in the grid (Outside of the scope of my internship/topic)

## 2 Look at the construction of custom loss functions for general machine learning models (XGBoost)

- Some physical laws can be encoded in the loss function
  - Loss:  $MSE + \Sigma(P) = 0?$
  - Prevents 0 predictions, still accounts for some part of the physics

## 3 Before learning the actual power factor, first check if we can accurately predict P

- How should we present the data?
- How do we ensure the model understands physical limits (E.g. Kirchhoff's current law)?
- Which features do we add?
  - Geographical data?
  - Weather data?
- Actual meat of the internship!

# A tale of two datasets

alllander

Original MSE data:

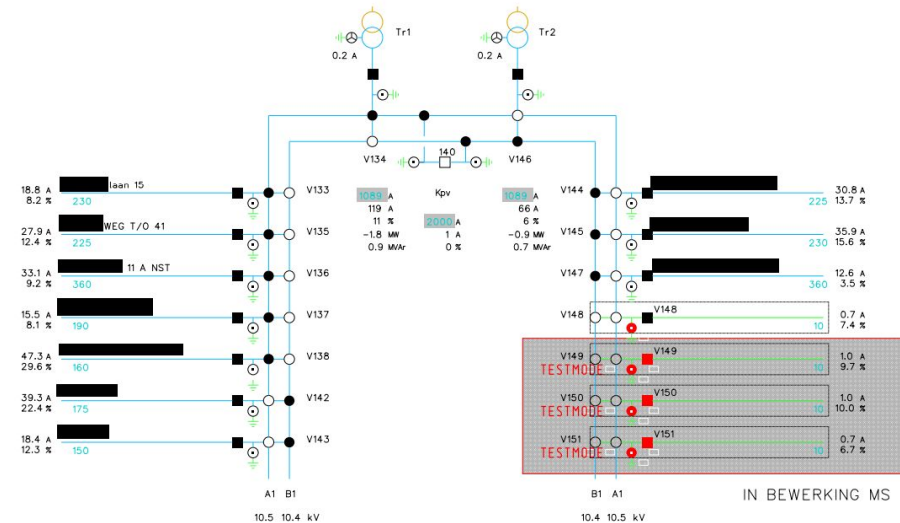
- 1 year, 10 stations
- Time series data (measurement / 5 min)
  - per station, per installation, per field
- Contains absolute current, voltage, active and reactive power
  - Unfortunately Kirchhoff's current law rarely held up (At worst 12 MW deviation, total ~35MW → ~ 30% deviation)

Decided to focus on Kirchhoff's current law as my primary physical law to incorporate

- How to combine spatial and temporal features?
- No good method found

Main features would be:

- Train: Current (I), voltage (U)
- Test: Active power (P), reactive power (Q)

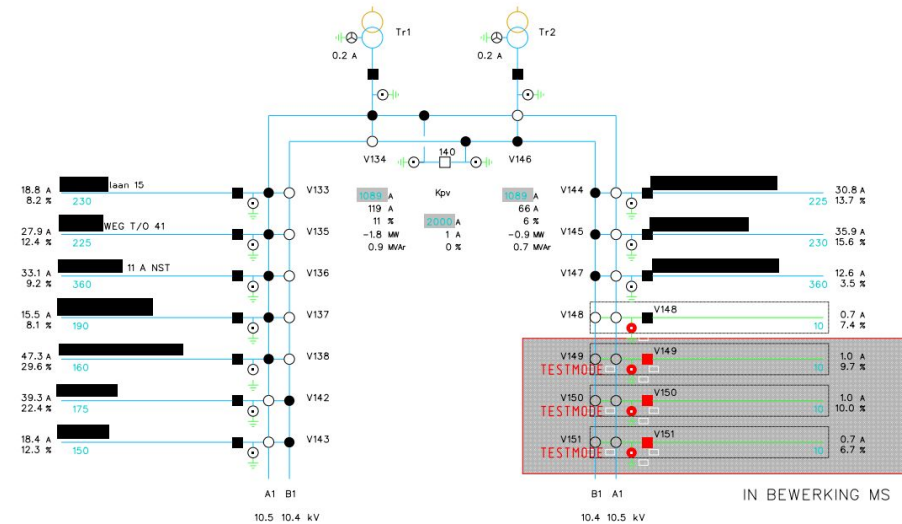


# A tale of two datasets

alllander

Decided to use two datasets

- Time series data
  - Time relation through lagging data:
    - P at t, t-1, t-2, etc
- Spatial data
  - No (real) time relation between data points
  - Single data point = entire station
  - Amount of features:  $F \times N$ 
    - $F$  = amount of fields per station
    - $N$  = amount of measurements per field
- Big problem!
  - Variable size of stations results in different amount of features
  - → Do we train per station or in general?
    - Resulting in idea for other experiment





	DATUM_TIJD	STATION	I_0	U_0	I_1	U_1	I_2	U_2	I_3	U_3	...	I_20	U_20	I_21	U_21	I_22	U_22
0	0.0	Tex	0.000000	0.0105	76.790000	0.0105	44.380000	0.0105	52.980000	0.0105	...	0.0	0.0000	0.00	0.0000	0.00	0.0000
1	0.0	Hby	2.080332	0.0105	25.878577	0.0105	2.743656	0.0105	9.843535	0.0105	...	0.0	0.0000	0.00	0.0000	0.00	0.0000
2	0.0	Dtn	65.649983	0.0105	45.709995	0.0105	6.767125	0.0105	0.000000	0.0000	...	0.0	0.0000	0.00	0.0000	0.00	0.0000
3	0.0	HFDP	0.000000	0.0210	0.000000	0.0210	127.780000	0.0210	126.570000	0.0210	...	0.0	0.0000	0.00	0.0000	0.00	0.0000
4	0.0	Grd	0.000000	0.0105	56.730000	0.0105	57.390000	0.0105	0.000000	0.0105	...	0.0	0.0105	22.19	0.0105	56.78	0.0105

Internship became more and more about data exploration than machine learning.

- What data to use (temporal vs spatial)?
- How to train (stations vs general)?
- Influence of different time of year?
- What/how to add knowledge about fields?

# Experiments

Machine learning models straight out of the box (scikit-learn)

- Gaussian process
- XGBoost


Experiment variables:

- Temporal or spatial data
- GP or XGBoost
- Train per station or on periods of 2 weeks
- Additional feature (Thanks to Jacco)
  - Treat all fields equal or differentiate between incoming and outgoing fields

Metrics used:

- MSE
- Accuracy on sign prediction

Difficult to determine when a prediction was good without visual inspection



# Results from experiments



Metric wise:

- Mixed results:
  - ~ 60-90% accuracy predicting sign (Promising)
    - Some outliers (50% or 95% accuracy)
  - ~ 5-10 MSE (painful)
    - Meaning an average error of  $\sim\sqrt{5}$  -  $\sqrt{10}$  MW

Notice a slight improvement in performance on training on individual stations.

However with visual inspection we soon see the results

- GP acc: 0.82, mse: 6.48
- XGB acc: 0.63, mse: 12.17



# Conclusion

## Disclaimers

- My code is still quite buggy/messy
  - Order of stations/fields does not match between spatial and temporal data
  - Splitting for cross validation between weeks and stations has inconsistencies
- Training for GPs had to be done on small sets due to intensive memory requirements
  - Therefore the training for XGBoost was limited to the same data size, to keep the comparison fair
- Error below obtained training GP on 2 weeks of data

**MemoryError:** Unable to allocate 75.7 GiB

## Take-away messages

- Further research required for data sanity
  - Does the data match the expected values?
  - Often I found myself in situations where it did not
- GPs are powerful, but perhaps unwieldy for a grid wide scope
  - Perhaps look in to (GP) models for individual stations?
- Further research in to developing a loss function
  - XGBoost seems widely used in Alliander
  - Incorporating simple constraints such as Kirchhoff's current law in to a loss function and implementing it
- Further research on the best way to represent the temporal and spatial information of the grid
  - Graph Neural Networks?
- Incorporating concrete constraints (Such as sum of predictions = 0) in to GPs is very complex
  - *"Incorporating Sum Constraints into Multitask Gaussian Processes"* Philipp Pilar